

CNN on Mnist Dataset

```
In [229... from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
import tensorflow as tf
from tensorflow.keras import utils
from tensorflow.keras.datasets import mnist
import seaborn as sns
from tensorflow.keras.initializers import RandomNormal
from tensorflow.python.keras import Input, Model
from tensorflow.keras.layers import Dense, Activation
from tensorflow.python.keras.layers import Dense, BatchNormalization
from tensorflow.keras.models import Sequential
%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
from keras.utils.np_utils import to_categorical

from tensorflow.keras import utils as np_utils
```

1.Data Loading

```
In [155... (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [156... # data is in 3d format
X_train.shape
```

Out[156... (60000, 28, 28)


```
In [157... print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)

Number of training examples : 10000 and each image is of shape (28, 28)

```
In [158... sample = 1
image = X_train[sample]

# plot the sample
fig = plt.figure
plt.imshow(image, cmap='gray')
plt.show()
```



```
In [159... img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
```

```
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)
```

```
In [160... x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.np_utils.to_categorical(y_train, num_classes)

y_test = keras.utils.np_utils.to_categorical(y_test, num_classes)

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

```
In [161... x_train.shape
```

```
Out[161... (60000, 28, 28, 1)
```

2.Architecture 1

```
In [162... batch_size = 128
num_classes = 10
epochs = 12
```

```
In [178... model1 = Sequential()

model1.add(Conv2D(20, kernel_size=(3, 3), activation='relu', input_shape=input_shape))

model1.add(Conv2D(filters=10, kernel_size = (3,3), activation="relu"))
model1.add(Conv2D(filters=20, kernel_size = (2,2), activation="relu"))

model1.add(MaxPooling2D(pool_size=(2,2)))
```

```

modell.add(Dropout(0.5))
modell.add(BatchNormalization())

modell.add(Conv2D(filters=5, kernel_size = (3,3), activation="relu"))
modell.add(Conv2D(filters=10, kernel_size = (2,2), activation="relu"))
modell.add(Dropout(0.5))
modell.add(Flatten())

modell.add(Dense(100,activation="relu"))

modell.add(Dense(10,activation="softmax"))

modell.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])

modell.summary()

```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
conv2d_23 (Conv2D)	(None, 26, 26, 20)	200
conv2d_24 (Conv2D)	(None, 24, 24, 10)	1810
conv2d_25 (Conv2D)	(None, 23, 23, 20)	820
max_pooling2d_5 (MaxPooling2D)	(None, 11, 11, 20)	0
dropout (Dropout)	(None, 11, 11, 20)	0
module_wrapper_12 (ModuleWrapper)	(None, 11, 11, 20)	80
conv2d_26 (Conv2D)	(None, 9, 9, 5)	905
conv2d_27 (Conv2D)	(None, 8, 8, 10)	210
dropout_1 (Dropout)	(None, 8, 8, 10)	0
flatten_4 (Flatten)	(None, 640)	0
module_wrapper_13 (ModuleWrapper)	(None, 100)	64100
module_wrapper_14 (ModuleWrapper)	(None, 10)	1010

Total params: 69,135
Trainable params: 69,095
Non-trainable params: 40

```
In [171... tf.config.run_functions_eagerly(True)
```

```
In [179... history = model1.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_te
```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/data/ops/dataset_ops.py:4212: UserWarning: Even though the `tf.config.experimental_run_functions_eagerly` option is set, this option does not apply to tf.data functions. To force eager execution of tf.data functions, please use `tf.data.experimental.enable_debug_mode()`.

"Even though the `tf.config.experimental_run_functions_eagerly` "

Epoch 1/12

469/469 [=====] - 101s 214ms/step - loss: 0.4657 - accuracy: 0.8501 - val_loss: 0.1199 - val
_accuracy: 0.9760

Epoch 2/12

469/469 [=====] - 101s 215ms/step - loss: 0.1425 - accuracy: 0.9562 - val_loss: 0.0548 - val
_accuracy: 0.9831

Epoch 3/12

469/469 [=====] - 100s 213ms/step - loss: 0.1075 - accuracy: 0.9667 - val_loss: 0.0448 - val
_accuracy: 0.9852

Epoch 4/12

469/469 [=====] - 100s 213ms/step - loss: 0.0947 - accuracy: 0.9705 - val_loss: 0.0472 - val
_accuracy: 0.9849

Epoch 5/12

469/469 [=====] - 100s 214ms/step - loss: 0.0836 - accuracy: 0.9741 - val_loss: 0.0350 - val
_accuracy: 0.9888

Epoch 6/12

469/469 [=====] - 101s 216ms/step - loss: 0.0752 - accuracy: 0.9765 - val_loss: 0.0392 - val
_accuracy: 0.9874

Epoch 7/12

469/469 [=====] - 100s 214ms/step - loss: 0.0682 - accuracy: 0.9790 - val_loss: 0.0323 - val
_accuracy: 0.9894

Epoch 8/12

469/469 [=====] - 99s 211ms/step - loss: 0.0624 - accuracy: 0.9803 - val_loss: 0.0307 - val
accuracy: 0.9908

Epoch 9/12

469/469 [=====] - 99s 212ms/step - loss: 0.0602 - accuracy: 0.9805 - val_loss: 0.0302 - val
accuracy: 0.9906

Epoch 10/12

469/469 [=====] - 100s 213ms/step - loss: 0.0537 - accuracy: 0.9824 - val_loss: 0.0325 - val
_accuracy: 0.9897

Epoch 11/12

```
469/469 [=====] - 100s 214ms/step - loss: 0.0551 - accuracy: 0.9822 - val_loss: 0.0290 - val_accuracy: 0.9908
Epoch 12/12
469/469 [=====] - 101s 215ms/step - loss: 0.0499 - accuracy: 0.9838 - val_loss: 0.0281 - val_accuracy: 0.9903
```

```
In [181... # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
In [182... score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,epochs+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_te

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
/usr/local/lib/python3.7/dist-packages/tensorflow/python/data/ops/dataset_ops.py:4212: UserWarning: Even though the `tf.config.experimental_run_functions_eagerly` option is set, this option does not apply to tf.data functions. To force eager execution of tf.data functions, please use `tf.data.experimental.enable_debug_mode()`.
"Even though the `tf.config.experimental_run_functions_eagerly` "
```

Test score: 0.02813357301056385

Test accuracy: 0.9902999997138977



3. Architecture 2

```
In [183... batch_size = 200
num_classes = 10
epochs = 10
```

```
In [192... model2 = Sequential()
model2.add(Conv2D(filters=5, kernel_size=3, strides=1, padding='same', activation='relu', input_shape=input_shape))
model2.add(MaxPooling2D(pool_size=2, strides=2, padding='same'))
model1.add(Dropout(0.5))
model2.add(Conv2D(filters=10, kernel_size=3, strides=1, padding='same', activation='relu'))
model2.add(MaxPooling2D(pool_size=2, strides=2, padding='same'))
model2.add(Conv2D(filters=5, kernel_size=3, strides=1, padding='same', activation='relu'))
model2.add(MaxPooling2D(pool_size=2, strides=2, padding='same'))
model1.add(BatchNormalization())
model1.add(Dropout(0.5))
```

```

model2.add(Flatten())
model2.add(Dense(100, activation='relu'))
model2.add(Dense(10, activation='softmax'))
model2.build()
model2.summary()

```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
conv2d_37 (Conv2D)	(None, 28, 28, 5)	50
max_pooling2d_15 (MaxPooling)	(None, 14, 14, 5)	0
conv2d_38 (Conv2D)	(None, 14, 14, 10)	460
max_pooling2d_16 (MaxPooling)	(None, 7, 7, 10)	0
conv2d_39 (Conv2D)	(None, 7, 7, 5)	455
max_pooling2d_17 (MaxPooling)	(None, 4, 4, 5)	0
flatten_8 (Flatten)	(None, 80)	0
module_wrapper_25 (ModuleWra	(None, 100)	8100
module_wrapper_26 (ModuleWra	(None, 10)	1010
Total params: 10,075		
Trainable params: 10,075		
Non-trainable params: 0		

```

In [195... model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model2.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_te

```

Epoch 1/10

1/300 [.....] - ETA: 37s - loss: 2.3208 - accuracy: 0.0850

/usr/local/lib/python3.7/dist-packages/tensorflow/python/data/ops/dataset_ops.py:4212: UserWarning: Even though the `tf.config.experimental_run_functions_eagerly` option is set, this option does not apply to tf.data functions. To force eager execution of tf.data functions, please use `tf.data.experimental.enable_debug_mode()`.

"Even though the `tf.config.experimental_run_functions_eagerly` "

300/300 [=====] - 37s 123ms/step - loss: 0.6301 - accuracy: 0.8112 - val_loss: 0.2513 - val_accuracy: 0.9239


```

Epoch 2/10
300/300 [=====] - 37s 122ms/step - loss: 0.2125 - accuracy: 0.9336 - val_loss: 0.1726 - val_
accuracy: 0.9463
Epoch 3/10
300/300 [=====] - 37s 122ms/step - loss: 0.1594 - accuracy: 0.9503 - val_loss: 0.1395 - val_
accuracy: 0.9555
Epoch 4/10
300/300 [=====] - 37s 122ms/step - loss: 0.1330 - accuracy: 0.9586 - val_loss: 0.1085 - val_
accuracy: 0.9659
Epoch 5/10
300/300 [=====] - 37s 122ms/step - loss: 0.1135 - accuracy: 0.9641 - val_loss: 0.1032 - val_
accuracy: 0.9677
Epoch 6/10
300/300 [=====] - 36s 120ms/step - loss: 0.0992 - accuracy: 0.9687 - val_loss: 0.0909 - val_
accuracy: 0.9718
Epoch 7/10
300/300 [=====] - 36s 121ms/step - loss: 0.0900 - accuracy: 0.9717 - val_loss: 0.0791 - val_
accuracy: 0.9735
Epoch 8/10
300/300 [=====] - 36s 121ms/step - loss: 0.0850 - accuracy: 0.9734 - val_loss: 0.0802 - val_
accuracy: 0.9743
Epoch 9/10
300/300 [=====] - 37s 123ms/step - loss: 0.0764 - accuracy: 0.9761 - val_loss: 0.0706 - val_
accuracy: 0.9766
Epoch 10/10
300/300 [=====] - 37s 124ms/step - loss: 0.0740 - accuracy: 0.9766 - val_loss: 0.0717 - val_
accuracy: 0.9751

```

In [196...

```

score = model2.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,epochs+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_te

# we will get val_loss and val_acc only when you pass the paramter validation_data

```


```
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/data/ops/dataset_ops.py:4212: UserWarning: Even though the `tf.config.experimental_run_functions_eagerly` option is set, this option does not apply to tf.data functions. To force eager execution of tf.data functions, please use `tf.data.experimental.enable_debug_mode()`.
 "Even though the `tf.config.experimental_run_functions_eagerly` "

Test score: 0.0717203095553055
 Test accuracy: 0.9750999808311462



4.Architecture 3

```
In [221... batch_size = 200
num_classes = 10
epochs = 15
```

```
In [222... model3 = Sequential()
model3.add(Conv2D(filters=5, kernel_size=3, strides=2, padding='same', activation='relu', input_shape=input_shape))
model3.add(MaxPooling2D(pool_size=2, strides=2, padding='same'))
model3.add(Dropout(0.5))
model3.add(Conv2D(filters=10, kernel_size=3, strides=2, padding='same', activation='relu'))
model3.add(MaxPooling2D(pool_size=2, strides=2, padding='same'))
model3.add(Dropout(0.5))
model3.add(Conv2D(filters=5, kernel_size=3, strides=2, padding='same', activation='relu'))
model3.add(MaxPooling2D(pool_size=2, strides=2, padding='same'))
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Flatten())
model3.add(BatchNormalization())
model3.add(Dropout(0.5))
model3.add(Dense(100, activation='relu'))
model3.add(Dense(10, activation='softmax'))
model3.build()
model3.summary()
```

Model: "sequential_34"

Layer (type)	Output Shape	Param #
conv2d_61 (Conv2D)	(None, 14, 14, 5)	50
max_pooling2d_36 (MaxPooling)	(None, 7, 7, 5)	0
dropout_24 (Dropout)	(None, 7, 7, 5)	0
conv2d_62 (Conv2D)	(None, 4, 4, 10)	460
max_pooling2d_37 (MaxPooling)	(None, 2, 2, 10)	0
dropout_25 (Dropout)	(None, 2, 2, 10)	0
conv2d_63 (Conv2D)	(None, 1, 1, 5)	455
max_pooling2d_38 (MaxPooling)	(None, 1, 1, 5)	0
module_wrapper_43 (ModuleWra	(None, 1, 1, 5)	20
dropout_26 (Dropout)	(None, 1, 1, 5)	0

flatten_12 (Flatten)	(None, 5)	0
module_wrapper_44 (ModuleWra	(None, 5)	20
dropout_27 (Dropout)	(None, 5)	0
module_wrapper_45 (ModuleWra	(None, 100)	600
module_wrapper_46 (ModuleWra	(None, 10)	1010
=====		
Total params: 2,615		
Trainable params: 2,595		
Non-trainable params: 20		

```
In [223... model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model3.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_te
```

```
Epoch 1/15
2/300 [.....] - ETA: 14s - loss: 2.4454 - accuracy: 0.0950
/usr/local/lib/python3.7/dist-packages/tensorflow/python/data/ops/dataset_ops.py:4212: UserWarning: Even though the `
tf.config.experimental_run_functions_eagerly` option is set, this option does not apply to tf.data functions. To forc
e eager execution of tf.data functions, please use `tf.data.experimental.enable_debug_mode()`.
"Even though the `tf.config.experimental_run_functions_eagerly` "
300/300 [=====] - 18s 59ms/step - loss: 2.2762 - accuracy: 0.1394 - val_loss: 2.1980 - val_a
ccuracy: 0.2962
Epoch 2/15
300/300 [=====] - 17s 57ms/step - loss: 2.1475 - accuracy: 0.1764 - val_loss: 1.8840 - val_a
ccuracy: 0.3405
Epoch 3/15
300/300 [=====] - 17s 57ms/step - loss: 2.0519 - accuracy: 0.1965 - val_loss: 1.7574 - val_a
ccuracy: 0.4054
Epoch 4/15
300/300 [=====] - 18s 58ms/step - loss: 1.9945 - accuracy: 0.2059 - val_loss: 1.6889 - val_a
ccuracy: 0.3725
Epoch 5/15
300/300 [=====] - 18s 58ms/step - loss: 1.9674 - accuracy: 0.2121 - val_loss: 1.6504 - val_a
ccuracy: 0.3524
Epoch 6/15
300/300 [=====] - 18s 59ms/step - loss: 1.9563 - accuracy: 0.2162 - val_loss: 1.6133 - val_a
ccuracy: 0.3571
Epoch 7/15
300/300 [=====] - 17s 57ms/step - loss: 1.9405 - accuracy: 0.2233 - val_loss: 1.6329 - val_a
ccuracy: 0.4147
Epoch 8/15
```

```

300/300 [=====] - 17s 57ms/step - loss: 1.9323 - accuracy: 0.2284 - val_loss: 1.6129 - val_a
ccuracy: 0.4404
Epoch 9/15
300/300 [=====] - 17s 56ms/step - loss: 1.9281 - accuracy: 0.2334 - val_loss: 1.6224 - val_a
ccuracy: 0.4171
Epoch 10/15
300/300 [=====] - 17s 57ms/step - loss: 1.9171 - accuracy: 0.2400 - val_loss: 1.6288 - val_a
ccuracy: 0.3966
Epoch 11/15
300/300 [=====] - 17s 57ms/step - loss: 1.9119 - accuracy: 0.2434 - val_loss: 1.6366 - val_a
ccuracy: 0.3338
Epoch 12/15
300/300 [=====] - 17s 57ms/step - loss: 1.9082 - accuracy: 0.2454 - val_loss: 1.6403 - val_a
ccuracy: 0.3305
Epoch 13/15
300/300 [=====] - 17s 57ms/step - loss: 1.9016 - accuracy: 0.2457 - val_loss: 1.6356 - val_a
ccuracy: 0.3324
Epoch 14/15
300/300 [=====] - 17s 58ms/step - loss: 1.8997 - accuracy: 0.2497 - val_loss: 1.6589 - val_a
ccuracy: 0.2780
Epoch 15/15
300/300 [=====] - 17s 57ms/step - loss: 1.8973 - accuracy: 0.2526 - val_loss: 1.6352 - val_a
ccuracy: 0.3115

```

```

In [225... score = model3.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,epochs+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_te

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

```

```
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/data/ops/dataset_ops.py:4212: UserWarning: Even though the `tf.config.experimental_run_functions_eagerly` option is set, this option does not apply to tf.data functions. To force eager execution of tf.data functions, please use `tf.data.experimental.enable_debug_mode()`.

"Even though the `tf.config.experimental_run_functions_eagerly` "

Test score: 1.6351990699768066

Test accuracy: 0.31150001287460327



5.Observation Table

```
In [230... from prettytable import PrettyTable
x = PrettyTable()

x = PrettyTable()
x.field_names=["Architecture Number", "epochs","Test Score", "Test Accuracy" ]

x.add_row(["Architecture 1", "12", "0.028","99" ])
```

```
x.add_row(["Architecture 2", "10", "0.071", "97" ])
x.add_row(["Architecture 3", "15", "1.63", "31" ])

print(x)
```

Architecture Number	epochs	Test Score	Test Accuracy
Architecture 1	12	0.028	99
Architecture 2	10	0.071	97
Architecture 3	15	1.63	31

6.Observations:

- 1.First two architecture gives good accuracy score.
- 2.Third architecture gives poor prformance